

Index 01

*“Ego kills knowledge,
as knowledge requires learning,
and learning requires humility”
- Rolsey*



RONIN™
ENGINEER

Outline

1. Query Optimization
2. Query Execution
 - SQL Execution
 - Execution Plan
3. Practices

1. Query Optimization

1.1. Query Optimization

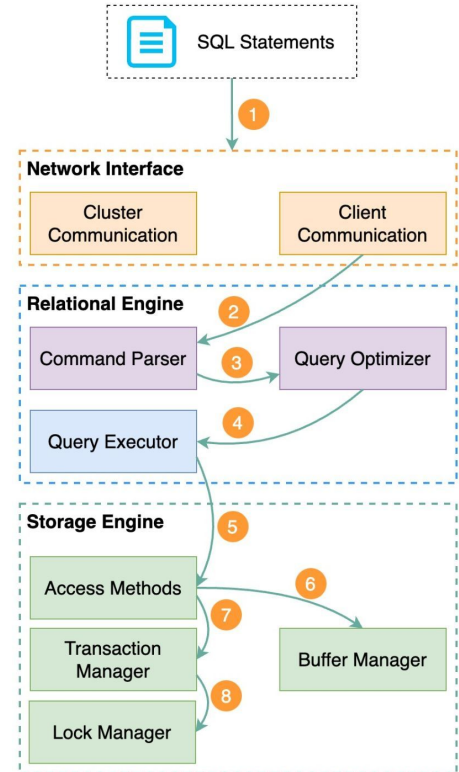
- There are two parts:
 - **Direct Query Optimization: changes to queries and indexes**
 - Rewrite query
 - Index
 - ...
 - **Indirect Query Optimization: changes to data and access patterns**
 - Changes to data: reducing the size of data, move old data to cold storage
 - Denormalization
 - Partitioning
 - Defragment (operation)
 - ...
- Optimize in the order: direct → indirect

2. Query Execution

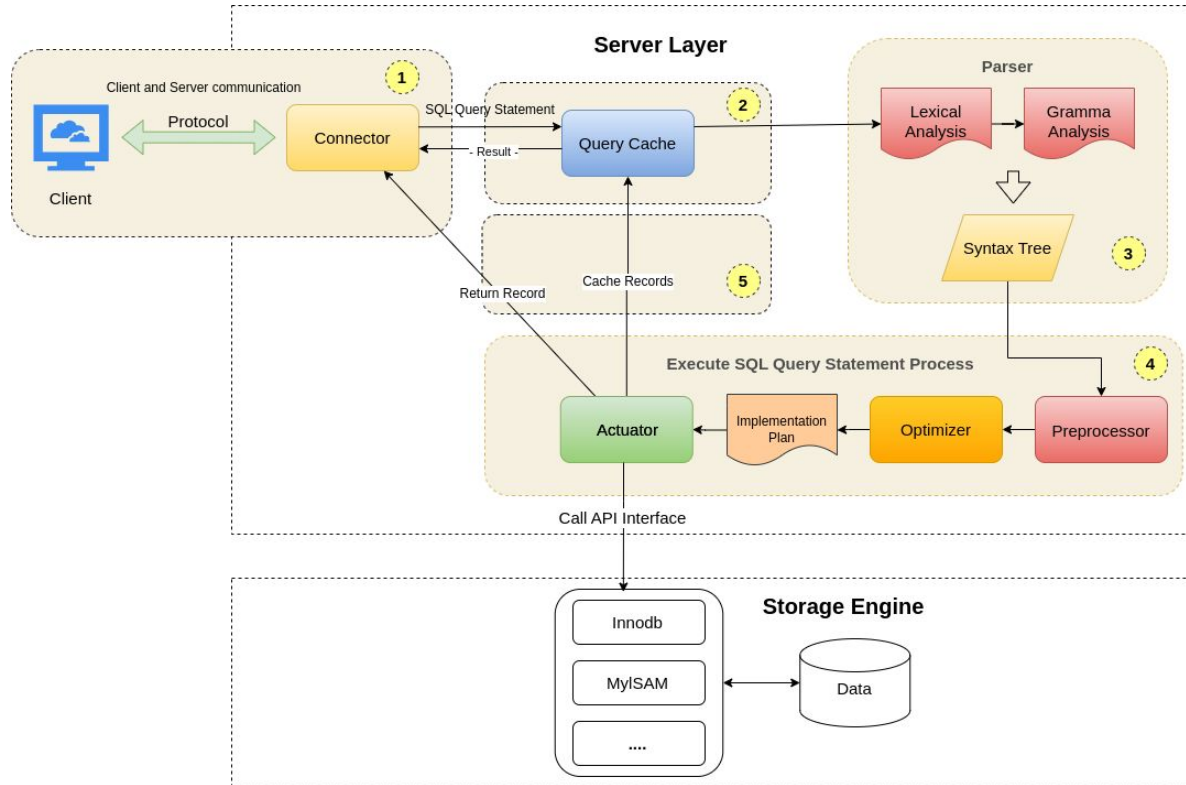
2.1. SQL Execution

2.1.1. How SQL Executed in DB?

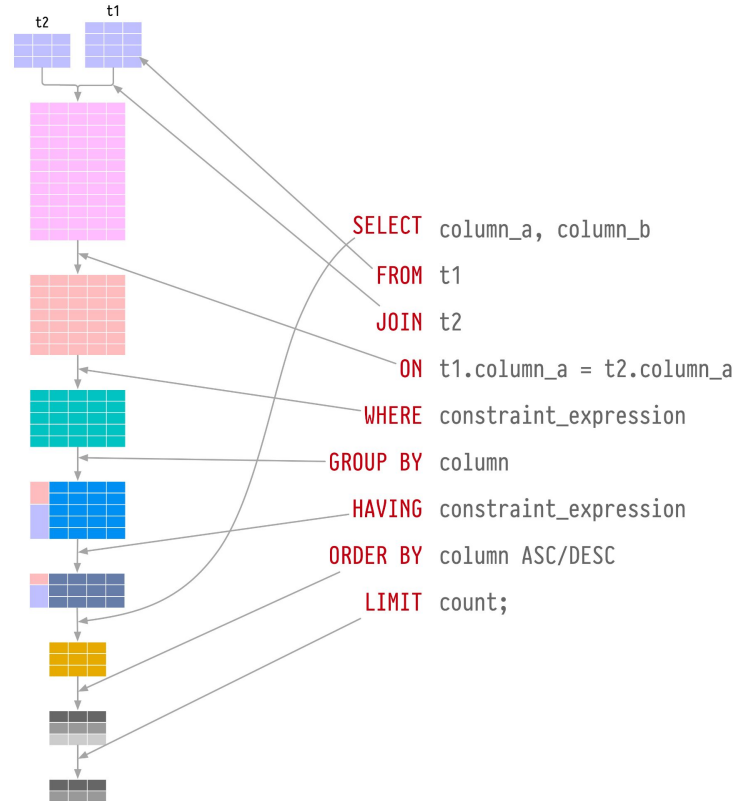
- Step 1: Client submit statement to RDBMS
- Step 2: RDBMS received SQL from Network Interface
- Step 3: The Parser preprocesses, parses SQL into a query tree
- Step 4: **Optimizer generates execution plans then choose one plan**
- Step 5: Executor retrieves data from storage engine based on the plan
- Step 6: If statement is read-only → Buffer Manager
- Step 7: If statement is a write → Transaction Manager
- Step 8: During a transaction, Lock manager ensure ACID properties



2.1.1. How SQL Executed in DB?



2.1.2. SQL Execution Order



2.2. Execution Plan

2.2.1. Execution Plan (Postgres)

- Execution plan is a detailed, step-by-step description of how RDBMS will execute a specific SQL query → **important to optimize query**
- Syntax:
 - EXPLAIN: to get basic information about the plan
 - ANALYZE: to get more concrete information about the plan
 - BUFFERS: information about cache hits/misses
 - FORMAT: (recommended) reformat output to YAML / JSON
 - Example: EXPLAIN (ANALYZE, BUFFERS, FORMAT YAML) SELECT ...

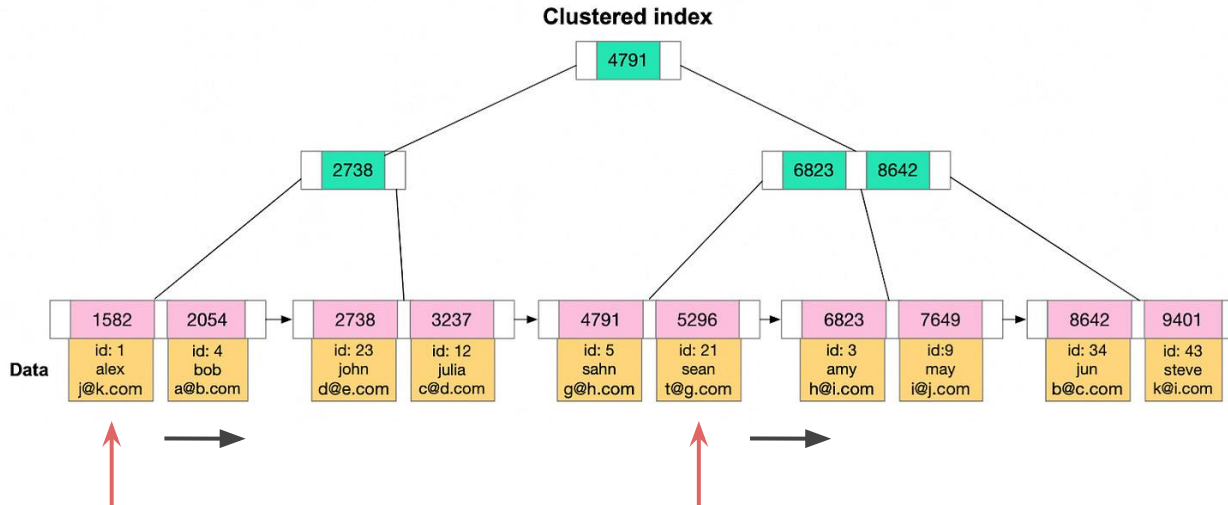
```
❏ QUERY PLAN
1  Index Scan using tickets_pkey on tickets  (cost=0.43..8.45 rows=1 width=104) (actual time=0.576..0.580 rows=1 loops=1)
2    Index Cond: (ticket_no = '0005434578291'::bpchar)
3 Planning Time: 0.098 ms
4 Execution Time: 0.622 ms
```

2.2.2. Types of Scanning

- Sequential Scan
- Index Scan
- Index Only Scan
- Bitmap Index Scan + Bitmap Heap Scan

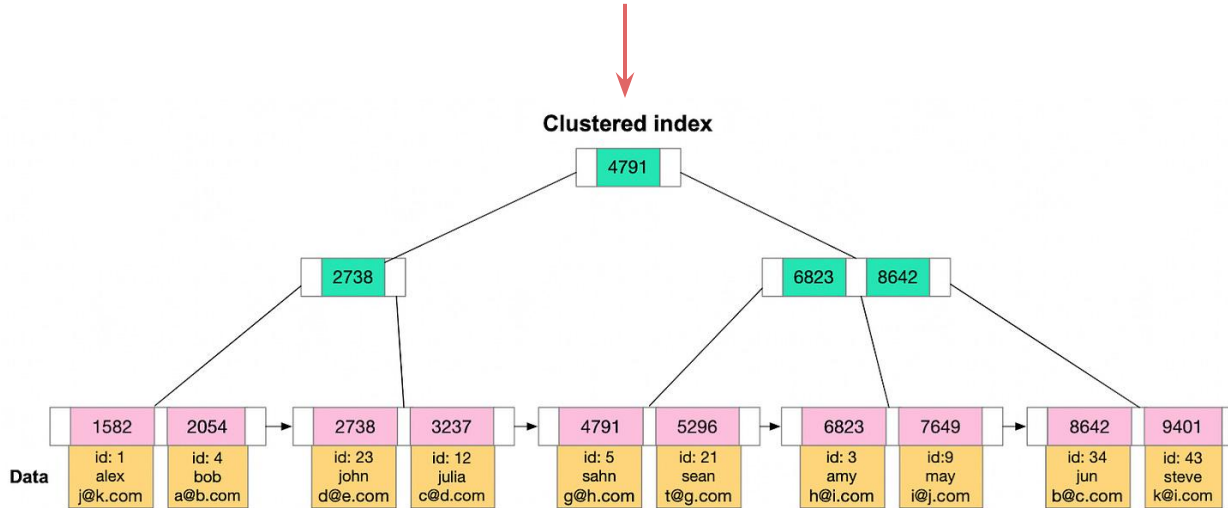
2.2.2. Types of Scanning

- **Sequential Scan:** scanning through the source table row-by-row without using any index
- **Parallel Sequential Scan:** using multi workers (threads) to scan



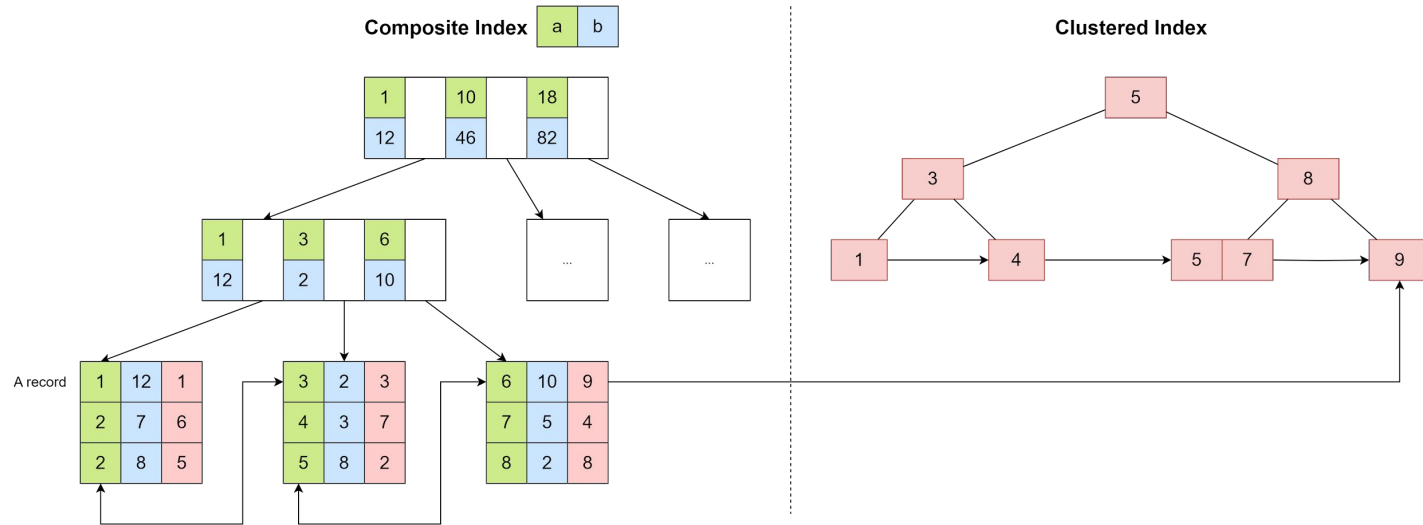
2.2.2. Types of Scanning

- **Index Scan:** using an index to determine which table rows match, and then retrieve the rows from the table



2.2.2. Types of Scanning

- **Index Only Scan:** retrieving the actual query result directly from the index, and avoids accessing the table itself for the requested data



2.2.2. Types of Scanning

- **Bitmap Index Scan + Bitmap Heap Scan:** using an index to generate a bitmap of which parts of the table likely contain the matching rows, and then access the actual table to get these rows using the bitmap - this is particularly useful to combine different indexes

2.3. Reading Execution Plan

- **Node Type:** types of operation
 - Seq Scan
 - Index Scan
 - Index Only Scan
 - Bitmap Index Scan
 - Nested Loop Join
 - Sort
 - Limit
 - ...
- **Index Name:** the index in used
- **Index Cond:** the condition used to scan on the index

```
explain (analyze, format yaml)
select * from tickets
where ticket_no = '0005434578291';

1  ∨ - Plan:
2      Node Type: "Index Scan"
3      Parallel Aware: false
4      Async Capable: false
5      Scan Direction: "Forward"
6      Index Name: "tickets_pkey"
7      Relation Name: "tickets"
8      Alias: "tickets"
9      Startup Cost: 0.43
10     Total Cost: 8.45
11     Plan Rows: 1
12     Plan Width: 104
13     Actual Startup Time: 0.250
14     Actual Total Time: 0.252
15     Actual Rows: 1
16     Actual Loops: 1
17     Index Cond: "(ticket_no = '0005434578291'"
18     Rows Removed by Index Recheck: 0
19     Planning Time: 0.775
20     Triggers:
21     Execution Time: 0.387
```

2.3. Reading Execution Plan

- Estimate Fields:
 - **Startup Cost:** the estimated amount of overhead necessary to start the operation (get the first record)
 - **Total Cost:** the estimated total cost of this operation and its descendants
 - **Plan Rows:** the number of rows the planner expects to be returned by the operation
 - **Plan Width:** the estimated average size of each row

```
explain (analyze, format yaml)
select * from tickets
where ticket_no = '0005434578291';

1  ∨ - Plan:
2      Node Type: "Index Scan"
3      Parallel Aware: false
4      Async Capable: false
5      Scan Direction: "Forward"
6      Index Name: "tickets_pkey"
7      Relation Name: "tickets"
8      Alias: "tickets"
9      Startup Cost: 0.43
10     Total Cost: 8.45
11     Plan Rows: 1
12     Plan Width: 104
13     Actual Startup Time: 0.250
14     Actual Total Time: 0.252
15     Actual Rows: 1
16     Actual Loops: 1
17     Index Cond: "(ticket_no = '0005434578291'"
18     Rows Removed by Index Recheck: 0
19     Planning Time: 0.775
20     Triggers:
21     Execution Time: 0.387
```

2.3. Reading Execution Plan

- Actual Value Fields:
 - **Actual Startup Time:** the amount of time it takes to get the first row out of the operation
 - **Actual Total Time:** the actual amount of time spent on this operation and all of its children
 - **Actual Rows:** the number of rows returned by the operation
 - **Actual Loops:** the number of times the operation is executed

```
explain (analyze, format yaml)
select * from tickets
where ticket_no = '0005434578291';

1  ∨ - Plan:
2      Node Type: "Index Scan"
3      Parallel Aware: false
4      Async Capable: false
5      Scan Direction: "Forward"
6      Index Name: "tickets_pkey"
7      Relation Name: "tickets"
8      Alias: "tickets"
9      Startup Cost: 0.43
10     Total Cost: 8.45
11     Plan Rows: 1
12     Plan Width: 104
13     Actual Startup Time: 0.250
14     Actual Total Time: 0.252
15     Actual Rows: 1
16     Actual Loops: 1
17     Index Cond: "(ticket_no = '0005434578291'"
18     Rows Removed by Index Recheck: 0
19     Planning Time: 0.775
20     Triggers:
21     Execution Time: 0.387
```

2.3. Reading Execution Plan

- Buffer fields:
 - **Shared Hit Blocks:** number of blocks read from cached indexes/tables

```
explain (analyze, format yaml)
select * from tickets
where ticket_no = '0005434578291';

1  ∨ - Plan:
2      Node Type: "Index Scan"
3      Parallel Aware: false
4      Async Capable: false
5      Scan Direction: "Forward"
6      Index Name: "tickets_pkey"
7      Relation Name: "tickets"
8      Alias: "tickets"
9      Startup Cost: 0.43
10     Total Cost: 8.45
11     Plan Rows: 1
12     Plan Width: 104
13     Actual Startup Time: 0.250
14     Actual Total Time: 0.252
15     Actual Rows: 1
16     Actual Loops: 1
17     Index Cond: "(ticket_no = '0005434578291'"
18     Rows Removed by Index Recheck: 0
19     Planning Time: 0.775
20     Triggers:
21     Execution Time: 0.387
```

3. Practices

3.1. Notes before practices

- More Datasets: [Kaggle](#)
- No way to show all candidate plans, they are removed right after optimizer plans

Demo

3.3. Key Takeaways

- Btree for range query, Hash for equal query.
- Index is suitable for fetch a small number of records.
- **Use composite indexes**, Ordering matters: high cardinality first.
- Limit the number of indexes by **leveraging index condition pushdown**.
- Leverage covering index (by using the INCLUDE for PostgreSQL).
- Do not rely on framework, lib to generate SQL.
- Inspect the execution plan of the generated SQL.
- PostgreSQL and MySQL do not work the same. Practice and practice with multiple databases
- Optimization in depth using pg_stat_statements.

<https://www.postgresql.org/docs/current/pgstatstatements.html>

Recap

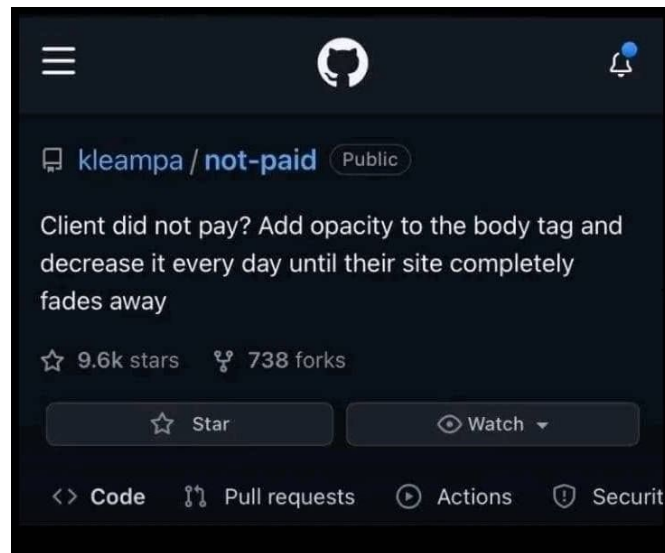
- Query Optimization includes 2 part: direct + indirect.
- Index is suitable for fetch a small result set.
- Inspect the execution plan of every production queries using EXPLAIN ANALYZE.

References

- <https://www.postgresqltutorial.com/postgresql-tutorial/postgresql-copy-table/>
- <https://blog.nodeswat.com/making-slow-queries-fast-with-composite-indexes-in-mysql-eb452a8d6e46>
- <https://www.pgmustard.com/blog/2019/9/17/postgres-execution-plans-field-glossary>
- <https://www.youtube.com/watch?v=LS-uE1V31IE>
- <https://www.pgmustard.com/blog>
- <https://www.percona.com/blog/mysql-101-how-to-find-and-tune-a-slow-sql-query/>

Homework

- Optimize 1 query of your project
 - Show solution
 - Explain why
- Index for the schema of booking flights



Thank you 🙏

